

COMPUTER LAB

Bioacoustics in Praat

Evolution of Language and Music

Goals Sounds are the raw materials in the study of language and music. In this lab we'll learn how to use Praat for analysing and editing sounds. We'll explore sound signals and look at how they relate to the things we perceive, such as words, melodies or rhythms.

1 Getting started

Praat is a free and open-source computer program widely used in phonetics (the study of human speech) and bioacoustics. It is a swiss-army knife containing many tools for visualising, analysing and synthesizing sounds.

EXERCISE 1

- Go to www.praat.org, find the download page for your favourite operating system and follow the installation instructions until you have started the Praat program (usually this involves double-clicking a beautiful pink icon).
- You will see two windows: **Praat objects** and **Praat picture**. **Praat objects** is where the sounds you are editing or analyzing will appear. **Praat picture** is where you can visualize the output of various analyses.

2 The anatomy of a sound

From the **Praat objects** window, navigate to *Open > Read from file*, or type Ctrl-O (Cmd-O on Mac). In the `materials` folder, you'll find a file called `sine.wav`. Open and load it into Praat. Now that we have a Praat object, let's have a look at what we can learn. First, let's play the sound.

EXERCISE 2

- Play the sound by selecting it from **Praat objects** and clicking **Play**.

Last updated on October 24, 2018

Soundwaves and spectrums The waveform is the most straightforward visual representation of a sound. The waveform is a plot of how the air pressure changes over time.

EXERCISE 3

- Click **View & Edit** to look at the waveform of our sound. You'll see two visual representations of the sound. The waveform is the upper one.
- In the **View & Edit** window, zoom in on the waveform until you can clearly see the shape of the sound waves.

You'll notice that this sound wave consists of a constantly repeating pattern. Each repetition of this pattern constitutes one vibration. The number of vibrations per second is called the *frequency* of the sound. Let's try to find out the frequency of the sound we've opened. To do this, we'll use a different representation of the sound, called a *Spectrum*. The Spectrum can be stored in a new Praat object (apart from sounds, Praat objects can also represent other information, such as the results of various sound-analyses).

EXERCISE 4

- In the Praat objects window, select **Analyse Spectrum** -, and click **To Spectrum...** Accept the default settings by clicking **OK**.
- Select the Spectrum object (if it isn't already selected) and visualize it by clicking the **View & Edit** button.
- Study the window and play around with it for a while: click anywhere in the plot, drag the mouse. What does the x-axis represent? What does the y-axis represent?
- Find the x-coordinate of the peak in the spectrum as precisely as possible. You'll probably need to zoom in a bit to do this accurately (tip: select the area around the peak you want to study and select "Zoom to selection" from the **View** menu at the top of the window, or press **Ctrl-n**). What is the frequency of the sound?

As you have heard, and seen, this sound is not particularly exciting. Let's look at a more interesting sound.

EXERCISE 5

- Load the file **bassoon.wav** into a Praat object
- Listen to both sounds (**bassoon.wav** and **sine.wav**) and compare. Do you hear any similarities, if so which? Which differences do you hear?
- Open the waveform view and zoom in (to somewhere in the middle of the sound) until you can see the individual vibrations of air-pressure (you can use the same zoom to selection technique that you used previously).

You should notice that the individual vibrations form a self-repeating pattern.

EXERCISE 6

- Find the shortest pattern in the waveform that contains no repetitions (Praat may already have marked this for you).
- Place the cursor at the start of the pattern, write down the exact time marking of the cursor.
- Place the cursor at the end of the pattern (exactly where it begins to repeat itself again), write down the exact time marking of the cursor.
- ? Using the two time markings, calculate the frequency (in repetitions per second) of the pattern you found.

The frequency you just found—the frequency of the shortest non-repeating pattern—is called the *fundamental frequency*. The fundamental frequency usually (but not always) corresponds to perceived pitch. As we will see now, sounds often contain many more frequencies, which can be discovered by looking at the spectrum.

EXERCISE 7

- ? Having analyzed the fundamental frequency of `bassoon.wav` and frequency of `sine.wav`, can you now, more precisely, describe the similarity between the two sounds?

EXERCISE 8

- Create a Spectrum object of `bassoon.wav` and display it with View & Edit.
- Can you find a peak in the spectrum corresponding to the frequency you found before?
- ★ Does the pitch that we perceive (the fundamental frequency) always correspond to the frequency of the highest peak in the spectrum?
- Read the frequencies of some other peaks in the spectrum. What do you notice about their relation to each other?

The peaks you found in the spectrum are called harmonics. The same note on various instruments may have the same pitch, but the energy distribution over the harmonics varies, resulting in different *timbres*. The same principle allows us to distinguish between different vowels.

2.1 The waveform and spectrogram

Now we'll look at human vocalizations.

EXERCISE 9

- Load the files `baby-1.wav` and `baby-2.wav` into Praat and listen to both sounds.^a
- Click View & Edit to look at the waveform for one of the files. Without zooming in, which properties of the sound can you recognize by just looking at the waveform?

As you can hear and see, these sounds are more complex than the sounds we've dealt with so far. The previous two sounds didn't change in pitch and maintained a (relatively) constant timbre throughout their duration. In the new sounds, the pattern of vibrations is continuously changing. Counting vibrations or looking at the spectrum will not be able to tell us much. With these sort of sounds, a *spectrogram* is a much more informative visualisation. You can view the spectrogram in the View & Edit window, just below the waveform. However, we're going to explore some Praat functionality to draw two spectrograms above each other in a picture.

We've seen how to view and edit Praat objects. Praat has different viewers for different objects. In these viewers, you can interact with the objects and zoom in to regions of interest. However, when you're, for example, writing a paper, you want to draw nice pictures containing these visualisations. For this reason, most Praat objects can be drawn into the **Praat picture** window. The praat picture, in turn, can be exported to various image formats.

EXERCISE 10

- Select one of the two baby sounds.
- In the *Praat picture* window draw a rectangle with a width of six and height of four (click and drag the mouse).
- Create a Spectrogram object. Click on the "Analyse spectrum -" button. From there, click on the "To spectrogram..." button and accept the default settings.
- Select Spectrogram object that you just created, click "Paint..." (under the "Draw -" button) and accept the default settings.
- Draw a second rectangle below the first one. Use the second rectangle to draw the Spectrogram of the other baby sound.
- ? Suppose you have heard the two sounds, and are now given these two spectrograms. Would you be able to figure which spectrogram belongs to which baby sound? If so, how?
- ★ What information does a spectrogram visualize? What do the x- and y-axes represent? What does the darkness of pixels mean?

A common analysis used for sounds is the F0 analysis, or fundamental frequency analysis. As we've seen above, the fundamental frequency generally corresponds to perceived pitch. We can use Praat to draw a *pitch contour*.

^aDuring the lecture, you heard cries from a French and a German baby. These were used in a study done by [Mampe2009](#). The recordings in this lab were recorded for a recent follow-up study done by [Wermke2016](#) comparing German and Chinese babies. Have a look at the studies if you're interested, in particular the first one! Both are included in this lab's materials.

EXERCISE 11

- Erase your Praat picture, by going to the Praat picture window, and clicking ‘Erase all’ under the ‘edit’ menu.

If you want, you can change the color and thickness of the drawn lines to make them stand out better. To do this

EXERCISE 12

- Open the “Pen” menu, and set the line width to 2.0 (by clicking on ‘Line width...’)
- In the same menu, change the color from black to something else. For example, red.

Now we’ll run the F0 analysis and draw the results.

EXERCISE 13

- Go to the Praat objects window.
 - Select the *sound* object you want to analyze.
 - Under ‘Analyse periodicity’, click ‘To pitch...’
 - Draw the created pitch object using the same method we used earlier.
- ? How do you think does Praat construct the pitch contour given a sound? Think of the manual analyses we did before. Describe the process informally, i.e., you don’t need to be very precise.

3 Speech

Although we’re all very good at producing and interpreting speech sounds, recognizing sounds in waveforms in spectrograms is much harder (even though they contain the same information!).

3.1 Phonemes

Phonemes are the basic components of speech. The word “slit”, for example consists of a fricative, a liquid, a vowel, and a plosive or stop (both so called “consonants”). Plosives are generated by completely stopping the airflow for a fraction of a section, resulting in complete silence.

EXERCISE 14

- Load the file `slit.wav`
- Take a look at the waveform and spectrogram and listen to the file

By looking carefully at the waveform and spectrogram, see if you can identify the individual phonemes making up the word. This may be harder than you expect.

EXERCISE 15

- To verify your identifications, extract each phoneme into a separate Praat object. Select the phoneme in the sound signal (you can either drag in the waveform or in the spectrogram), and click “File” > “Extract selected sound (preserve times)”. This will create a new Praat object, untitled. Use the rename button to rename it s, l, i or t to help you remember which phoneme it contains.
- Create a spectrum (not a spectrogram) object for the s (fricative) and i (vowel) sound and compare the two
- Now compare the s and i spectrums to the corresponding part of the spectrogram for slit.

Previously, we looked at harmonic frequencies in the bassoon sound. Amplified harmonics in speech sounds show up as peaks in the spectrum, or dark spots in the spectrogram. These peaks are called formants. Vowels can be differentiated by looking at how their formants are distributed.

EXERCISE 16

- ? How can you identify a fricative in the spectrogram?
- ★ How can you recognize a plosive in the spectrogram? And in the waveform?

3.2 The sound of silence

Very small changes to the signal can sometimes have dramatic effects on perception. For example, inserting a small period of silence (silent interval) at specific places in words can create the effect of hearing an extra phoneme. In this final part of the lab we’ll explore the effect of inserting a small silence in our recording of “slit” at just the right place.

First, we’ll create a small silence to be inserted into the `slit.wav` sound. To find out an appropriate duration for this silence, we’ll look at a paper that investigated the effect of a silent interval in the word “slit”. Have a look at the methods section, as well as the graph with results, in the paper by **Marcus1978** that’s attached to this lab (`paper.pdf`). Use the graph summarizing their results to find a good duration for the silent interval.

EXERCISE 17

- In the Praat objects window, go to the menu “New” > “Sound” and click “Create sound from formula”
- Change the value of the “Name” field to “silence”.
- Adjust the end time to the duration of the silent interval that you found
- In the formula field, type “0” (zero)
- Click OK
- Open the View & Edit screen for your new sound
- Select the entire sound (have a look at the Select menu if you run into

issues)

- Copy it, using “Edit > Copy selection to Sound clipboard” or Ctrl-c

Now we’re going to insert the silence into our recording of the word “slit”.

EXERCISE 18

- Go to the View & Edit window for the sound `slit.wav`
- Using the spectrogram and waveform, find a spot in between the “s” and the “l” sound and place the cursor there

To prevent sudden jumps in the waveform, we should insert our silence at a moment where the wave crosses the zero line.

EXERCISE 19

- After having placed the cursor between the “s” and “l” sound, click on “Select” > “Move cursor to nearest zero crossing”
- Now insert the silence we copied earlier by clicking “Edit” > “Paste after selection”, or by pressing Ctrl-v.
- Play the sound. Which word do you hear?